`learn2learn`: A Library for Meta-Learning Research

Sébastien M. R. Arnold*[1], Praateek Mahajan[2], Debajyoti Datta[3], Ian Bunner[4], and Konstantinos Saitas Zarkias[5]

1 University of Southern California **2** Iterable, Inc. **3** University of Virginia **3** University of Waterloo **5** KTH Royal Institute of Technology and RISE - Research Institutes of Sweden, SICS

Abstract

Meta-learning researchers face two fundamental issues in their empirical work: prototyping and reproducibility. Researchers are prone to make mistakes when prototyping new algorithms and tasks because modern meta-learning methods rely on unconventional functionalities of machine learning frameworks. In turn, reproducing existing results becomes a tedious endeavour – a situation exacerbated by the lack of standardized implementations and benchmarks. As a result, researchers spend inordinate amounts of time on implementing software rather than understanding and developing new ideas.

This manuscript introduces `learn2learn`, a library for meta-learning research focused on solving those prototyping and reproducibility issues. `learn2learn` provides low-level routines common across a wide-range of meta-learning techniques (e.g. meta-descent, meta-reinforcement learning, few-shot learning), and builds standardized interfaces to algorithms and benchmarks on top of them. In releasing `learn2learn` under a free and open source license, we hope to foster a community around standardized software for meta-learning research.

Introduction

Meta-learning is the subfield of machine learning that endows computer programs with the ability of learning to learn. That is, the computer not only learns a behavior, but also how to adapt its behavior. To illustrate this difference, let us make an anology with the world of athleticism. If a learning program is akin to an athlete, a meta-learning program corresponds to the athlete-coach pair: it simultaneously learns a skill and how to teach it best. Meta-learning is appealing whenever the athlete is required to excel in multiple sports, as the coach can leverage shared aspects of each sport to accelerate mastery. Recent meta-learning methods emerged from this natural ability to multitask – e.g. meta-descent in optimization, meta-reinforcement learning in reinforcement learning, and few-shot learning when labelled data is limited – reaching state-of-the-art performance on vision, language, and robotic domains. (Sutton 1992; Duan, Schulman, et al. 2016; Miller, Matsakis, and Viola 2000; Lee et al. 2019; Brown et al. 2020; Metz et al. 2019; Nagabandi et al. 2018)

---

*Correspondance to seb.arnold@usc.edu

Unfortunately, modern meta-learning research is often slowed down by prototyping and reproducibility challenges. Prototyping algorithms is an error-prone process since meta-learning algorithms rely on supported but exotic functionalities of machine learning software. (e.g. gradient of optimization steps) Consequently, a slow prototyping phase reduces the number of ideas one can try and retain. It also directly impacts reproducibility: researchers are more likely to make mistake with someone else's idea than with their own. Combined, those pesky issues prevent meaningful comparisons across publications, ultimately reducing the impact of any publication in the field.

Why do we suffer prototyping and reproducibility issues?

Although a complete answer is outside the scope of this manuscript, we blame the lack of specialized software as a major culprit. A software library with specialized subroutines would reduce gaffes when prototyping, improving reproducibility too. A widely-adopted library also promotes standardized implementation of existing methods and benchmarks, an issue in prior research; for example, the community lost the original mini-ImageNet data splits from Vinyals et al. (2016), leaving subsequent work to replicate them as best they could. In summary, there is a need for a specialized software library which would alleviate many of the issues currently plaguing meta-learning research.

We introduce `learn2learn`, a software library that directly addresses prototyping and re-productibility issues in meta-learning research. `learn2learn` provides researchers with a unified and extensible interface to existing benchmarks, and a set of well-tested subroutines frequently used to implement meta-learning algorithms. The library also packages numerous examples replicating published methods, which can easily be adapted for comparisons on new benchmarks. `learn2learn` is implemented in Python to maintain compatiblity with the greater machine learning ecosystem. It extends PyTorch (Paszke et al. 2019) and leverages its fast linear algebra and automatic differentiation capabilites, while resorting to Cython (Behnel et al. 2011) when speed is required for data-handling. `learn2learn` is already used in our day-to-day research, and we hope its development will continue to benefit the wider meta-learning community.

The remainder of this document overviews the prototyping and reproducibility capabilities of `learn2learn`, including a review of related work.

Prototyping

Prototyping is essential in letting researchers quickly try new ideas. The faster the pro-totyping phase, the faster an idea can be retained (or discarded) for further exploration. `learn2learn` provides tools to accelerate two aspects of the prototyping phase: algorithms and domains.

Algorithms

Implementing meta-learning algorithms can be tricky. For example, many methods rely

on computing gradients of algorithms – rather than gradients of functions – which, while possible with modern machine learning frameworks (e.g. PyTorch, TensorFlow, JAX), is strenuous and prone to errors. (Finn, Abbeel, and Levine 2017; Jacobsen et al. 2019; Xu, Hasselt, and Silver 2018) However, this doesn't need be the case: many differentiable algorithms can be implemented with minor changes when given the right abstractions.[1]

```python
1  learned_update = l2l.optim.ParameterUpdate(
2      model.parameters(),
3      l2l.optim.KroneckerTransform(l2l.nn.KroneckerLinear)
4  )
5  clone = l2l.clone_module(model)   # torch.clone() for nn.Modules
6  updates = learned_update(   # similar API as torch.autograd.grad
7      loss(clone(X), y),
8      clone.parameters(),
9      create_graph=True,
10 )
11 # in-place, differentiable update of clone parameters
12 l2l.update_module(clone, updates)
13 # gradients w.r.t model's and learned_update's parameters
14 loss(clone(X), y).backward()
```

Snippet 1: Demonstration of differentiable optimization routines. Lines 1-4 instantiate a parameterized update function, which computes the gradient of a loss w.r.t. some module's parameters (clone, line 8) and passes them through a *gradient transform* – a module mapping gradients to updates. (here a KroneckerLinear, line 3) Line 5 creates a differentiable copy of the model, and line 11 updates that copy in-place such that the update is itself differentiable. Finally, line 13 backpropagates through the loss of the updated differentiable copy, thus computing gradients w.r.t. the "pre-update" model parameters and the KroneckerLinear parameters.

To ease the implementation of such methods, learn2learn exposes low-level routines for differentiable optimization in learn2learn.optim. These routines are tightly built around PyTorch's automatic differentiation engine, so as to maintain compatibility and extensibility. They can be used to express optimization algorithms such that their computational graph remains differentiable. Snippet 1 provides an example, which implements the 1-step learning loop of the linear Kronecker-factored optimizer from Arnold, Iqbal, and Sha (2019). (Note how both model and optimizer parameters are meta-learned; implementing the same functionality with vanilla PyTorch requires 10x the lines of code.) We've used those general-purpose routines to implement algorithms from the literatures of few-shot, meta-descent, and meta-reinforcement learning – including MAML (Finn, Abbeel, and Levine 2017), Hypergradient descent (Baydin et al. 2017), or ProMP (Rothfuss et al. 2018) among others.

Domains

Researchers have to design new domains – such as datasets, tasks, or environments – to develop and test new abilities of their programs. We refer to this other aspect of

---

[1]See for example Agrawal et al. (2019).

the research lifecycle as *prototyping new domains*. `learn2learn` can help prototype new domains for few-shot and meta-reinforcement learning.

```
1  dataset = l2l.data.MetaDataset(MyDataset())  # PyTorch dataset
2  transforms = [  # easy to define custom task transforms
3      l2l.data.transforms.NWays(dataset, n=5),
4      l2l.data.transforms.KShots(dataset, k=1),
5      l2l.data.transforms.LoadData(dataset),
6      lambda task: [(random_rotation(x), y) for x, y in task]
7  ]
8  taskset = l2l.data.TaskDataset(dataset,
9                                 transforms,
10                                num_tasks=20000)
11 random_task = taskset.sample()  # sample one task
12 for task in taskset:  # iterate over all tasks
13     X, y = task
```

Snippet 2: The interface to `TaskDataset` and `TaskTransform`. Line 1 wraps an arbitrary PyTorch dataset with the `MetaDataset` class. Lines 2-7 define `TaskTransforms`: here, 5-ways 1-shot classification with a custom random-rotation task augmentation applied to each (x, y) pair. Lines 8-10 instantiate the `TaskDataset`, from which tasks can be sampled (line 11) or enumerated (lines 12-13).

For few-shot meta-learning, `learn2learn` provides a general `TaskDataset` class enabling sampling of smaller tasks in `learn2learn.data`. Those tasks are constructed through a series of `TaskTransforms`, which iteratively refine the description of the data in the task. Writing a new task transform is as easy as writing a Python function, but those functions can be made arbitrarily complex thanks to Python's callable objects. See Snippet 2 for an example. With the combination of `TaskDataset` and `TaskTransforms`, researchers can quickly develop fast custom data and task sampling schemes, while retaining compatibility with any PyTorch dataset; this lets them quickly iterate over ideas with small datasets and scale up to larger experiments with the same codebase.

```
1  def make_env():
2      env = l2l.gym.HalfCheetahForwardBackwardEnv()
3      return cherry.envs.StateNormalizer(env)
4
5  # use 16 processes, compatible with gym API
6  env = l2l.gym.AsyncVectorEnv([make_env for i in range(16)])
7  tasks = env.sample_tasks(20)
8  env.set_task(tasks[0])  # all processes run task 0
```

Snippet 3: Utilities for meta-reinforcement learning environments. Lines 1-3 instantiate the half-cheetah environment, with tasks defined as running forward or backward. This environment is then wrapped by cherry, an external reinforcement learning library. Line 6 forks 16 asynchronous workers, each with its own copy of the half-cheetah environment. Finally, lines 7-8 sample 20 tasks and assign the first one to all workers.

For meta-reinforcement learning, `learn2learn` provides a high-level `MetaEnv` interface in `learn2learn.gym` which can be used to bootstrap the design of OpenAI Gym en-

vironments. (Brockman et al. 2016) Environments that adhere to this interface can take advantage of specially designed utilities included in `learn2learn`; for example, the `AsyncVectorEnv` wrapper parallelizes the collection of episodes across multiple processes. (c.f. Snippet 3) Such environments also retain the Gym API, making them compatible with all popular reinforcement learning libraries. (Dhariwal et al. 2017; Duan, Chen, et al. 2016; Liang et al. 2017) Naturally, they also become compatible with the various meta-reinforcement learning algorithms implemented within the library.

The core prototyping tools toured in the above paragraphs open the door to more advanced developments such as online, incremental, or lifelong meta-learning. While this manuscript can only present a bird's-eye view of those tools, we invite the reader to the library's website and documentation for such advanced applications.

Reproducibility

The field of meta-learning is advancing rapidly, but progress is plagued by reproducibility issues. Those issues are often subtle and hard to spot. In meta-reinforcement learning for example, different papers have used different reward functions with the same environment, resulting in confusing comparisons: are the observed improvements due to algorithmic advancements or to changes in the reward function? To combat those insiduous issues, `learn2learn` includes a set of high-quality implementations for various meta-learning algorithms as well as standardized benchmarks for few-shot and meta-reinforcement learning.

```
1  meta_sgd = l2l.algorithms.GBML(
2      model,
3      l2l.optim.ModuleTransform(l2l.nn.Scale),
4  )
5  meta_curvature = l2l.algorithms.GBML(
6      model,
7      l2l.optim.MetaCurvatureTransform,
8  )
9  meta_kfo = l2l.algorithms.GBML(
10     model,
11     l2l.optim.KroneckerTransform(l2l.nn.KroneckerLinear),
12     adapt_transform=True,
13 )
```

Snippet 4: Implementation of Meta-SGD (lines 1-4), Meta-Curvature (lines 5-8), and Meta-KFO (lines 9-13) with the `GBML` wrapper. Each variant differs in how the fast-adaptation gradients are transformed, which is reflected through the `transform` and `adapt_transform` arguments.

Implementations

`learn2learn` provides high-level implementations for popular algorithms. These implementations build on top of the low-level routines from the previous section, and are thoroughly tested to replicate published works. They typically wrap around PyTorch modules to extend

them with specific meta-learning functionalities. For example, the `LearnableOptimizer` retains the familiar PyTorch `Optimizer` interface and extends it to learn arbitrary meta-optimization updates; similarly, the `GBML` augments PyTorch `Modules` to support fast-adaptation routines for few-shot and meta-reinforcement learning. `GBML` implementations of Meta-SGD, Meta-Curvature, Meta-KFO are available in Snippet 4. (Li et al. 2017; Park and Oliva 2019; Arnold, Iqbal, and Sha 2019)

With those high-level implementations and the standardized benchmark interface described below, we supply examples that exactly reproduce published experiments. Those examples serve three purposes. First, they validate the correctness of our implementation and a publication's claims; second, they illustrate how to use the library; third, they fill the need for unified standardized reproductions. As a by-product, those examples can be used to bootstrap further experimentation and analysis around a method of interest. For example, we could easily complement ANIL's (Raghu et al. 2019) original results on the Omniglot (Lake, Salakhutdinov, and Tenenbaum 2015) and mini-Imagenet (Vinyals et al. 2016) datasets with new results on CIFAR-FS (Bertinetto et al. 2018) and FC100 (Oreshkin, Rodríguez López, and Lacoste 2018).

```python
from learn2learn.vision import benchmarks
print(benchmarks.list_tasksets())
# ['omniglot', 'cifar-fs', 'fc100', 'mini-imagenet', ...]
tasksets = benchmarks.get_tasksets(  # standardized pipeline
    name='mini-imagenet',
    train_samples=10,
    train_ways=5)
task = tasksets.train.sample() # tasksets.train is a task dataset
```

Snippet 5: High-level API to standardized computer vision benchmarks. Line 2 prints the list of available tasksets. On lines 4-7, we instantiate the 5-ways 5-shots mini-ImageNet benchmark, which returns the `tasksets namedtuple` with `train`, `validation`, and `test` keys. Line 8 samples a task from the train `TaskDataset`.

Benchmarks

We use the low-level domains API to implement standard benchmarks in few-shot and meta-reinforcement learning settings. For few-shot learning, `learn2learn` provides classes to download and preprocess datasets commonly used by the community in `learn2learn.vision`. In addition, it also includes task definitions with the proper task-processing stages (image normalization, rotation, cropping) for commonly reported settings such as 5-ways 1-shot, 5-ways 5-shots, and 20-ways 5-shots. An example of those task definitions is described in Snippet 5.

The meta-reinforcement learning environments range from simple 2D-particle navigation to robotics control. In particular, we include simple to use wrappers for the recently proposed MetaWorld, a set of 50 gripper manipulation tasks with varying levels of difficulty. Those benchmark implementations should greatly simplify the replication and comparison of new

methods on well-studied settings. They, and other meta-reinforcement learning utilities, are included in `learn2learn.gym`.

Combined with the provided implementations, we hope `learn2learn` enables researchers to accelerate the process of correctly comparing their ideas against existing methods.

Related Work

Disparate implementations of individual algorithms set aside, two recent libraries tackle similar challenges as `learn2learn`.

The first one, `higher` (Grefenstette et al. 2019), aims to facilitate the implementation of "generalized inner-loop meta-algorithms" – in other words, the implementation of differentiable optimization algorithms. `higher` treats a model definition (e.g. a neural network) as a symbolic computational graph, for which they use one set of parameters or another based on user-specified context. This "stateless" parameterization is as expressive as the `learn2learn.optim` submodule – after all, both are implemented on top of PyTorch – but requires researchers to carefully understand when they are working with symbolic or declarative parts of their computation. Instead, `learn2learn` sticks with a stateful and declarative style, already familiar to the PyTorch research userbase. Moreover, `higher` completely forgoes reproducibility issues as its focus is on the implementation of novel algorithms.

The second one, `Torchmeta` (Deleu et al. 2019), intends to provide a unified interface to popular datasets, including classes to easily download and process them. Specifically, it focuses on standardized few-shot computer vision tasks, allowing researchers to easily swap one dataset for another. However, supporting new datasets with `Torchmeta` requires implementing a bridging class, even if the dataset is already in standard PyTorch format. On the other hand, `learn2learn`'s TaskDataset explicitly avoids such bridging classes, and is designed to be compatible with any PyTorch dataset (including text, speech, and others). `Torchmeta` also povides a thin algorithmic wrapper to demonstrate the usage of their library with gradient-based meta-learning algorithms; but, this wrapper is not compatible with the majority of PyTorch's layers, nor custom modules implemented by researchers. In comparison, `learn2learn`'s differentiable optimization submodule uniformly handles all PyTorch Modules (including custom ones), making it a more applicable research tool.

Overall, `learn2learn` offers a more general solution to the prototyping and reproduciblity issues encountered in day-to-day research. For example, at the time of this writing, neither of the above libraries supports meta-descent or meta-reinforcement learning.

Conclusion

This manuscript introduces `learn2learn`, a library that tackles prototyping and reproducibility issues in modern meta-learning. `learn2learn`'s low-level routines facilitate rapid

prototyping of new algorithms and domains. The library builds on those low-level routines to provide high-level implementations and standardized benchmarks API, which enable researchers to easily and faithfully compare different methods under different settings. Notably, both low- and high-level utilities are engineered to be general: they are compatible with a wide range of meta-learning techniques in few-shot learning, meta-reinforcement learning, or meta-optimization. Finally, `learn2learn` is released under the free and open source MIT licence, and the focus of continued development.

Acknowledgements

References

Agrawal, A., B. Amos, S. Barratt, S. Boyd, S. Diamond, and Z. Kolter. 2019. "Differentiable Convex Optimization Layers." In *Advances in Neural Information Processing Systems*.

Arnold, Sébastien M R, Shariq Iqbal, and Fei Sha. 2019. "When MAML Can Adapt Fast and How to Assist When It Cannot," October.

Baydin, Atilim Gunes, Robert Cornish, David Martinez Rubio, Mark Schmidt, and Frank Wood. 2017. "Online Learning Rate Adaptation with Hypergradient Descent," March.

Behnel, Stefan, Robert Bradshaw, Craig Citro, Lisandro Dalcin, Dag Sverre Seljebotn, and Kurt Smith. 2011. "Cython: The Best of Both Worlds." *Computing in Science & Engineering* 13 (2). IEEE: 31–39.

Bertinetto, Luca, Joao F Henriques, Philip Torr, and Andrea Vedaldi. 2018. "Meta-Learning with Differentiable Closed-Form Solvers."

Brockman, Greg, Vicki Cheung, Ludwig Pettersson, Jonas Schneider, John Schulman, Jie Tang, and Wojciech Zaremba. 2016. "OpenAI Gym," June.

Brown, Tom B, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, et al. 2020. "Language Models Are Few-Shot Learners," May.

Deleu, Tristan, Tobias Würfl, Mandana Samiei, Joseph Paul Cohen, and Yoshua Bengio. 2019. "Torchmeta: A Meta-Learning Library for PyTorch," September.

Dhariwal, Prafulla, Christopher Hesse, Oleg Klimov, Alex Nichol, Matthias Plappert, Alec Radford, John Schulman, Szymon Sidor, Yuhuai Wu, and Peter Zhokhov. 2017. "OpenAI Baselines." *GitHub Repository*.

Duan, Yan, Xi Chen, Rein Houthooft, John Schulman, and Pieter Abbeel. 2016. "Benchmarking Deep Reinforcement Learning for Continuous Control," April.

Duan, Yan, John Schulman, Xi Chen, Peter L Bartlett, Ilya Sutskever, and Pieter Abbeel. 2016. "RL^2: Fast Reinforcement Learning via Slow Reinforcement Learning," November.

Finn, Chelsea, Pieter Abbeel, and Sergey Levine. 2017. "Model-Agnostic Meta-Learning for Fast Adaptation of Deep Networks." *arXiv Preprint arXiv:1703. 03400.*

Grefenstette, Edward, Brandon Amos, Denis Yarats, Phu Mon Htut, Artem Molchanov, Franziska Meier, Douwe Kiela, Kyunghyun Cho, and Soumith Chintala. 2019. "Generalized Inner Loop Meta-Learning," October.

Jacobsen, Andrew, Matthew Schlegel, Cameron Linke, Thomas Degris, Adam White, and Martha White. 2019. "Meta-Descent for Online, Continual Prediction," July.

Lake, Brenden M, Ruslan Salakhutdinov, and Joshua B Tenenbaum. 2015. "Human-Level

Concept Learning Through Probabilistic Program Induction." *Science* 350 (6266): 1332–8.

Lee, Kwonjoon, Subhransu Maji, Avinash Ravichandran, and Stefano Soatto. 2019. "Meta-Learning with Differentiable Convex Optimization," April.

Li, Zhenguo, Fengwei Zhou, Fei Chen, and Hang Li. 2017. "Meta-SGD: Learning to Learn Quickly for Few-Shot Learning," July.

Liang, Eric, Richard Liaw, Philipp Moritz, Robert Nishihara, Roy Fox, Ken Goldberg, Joseph E Gonzalez, Michael I Jordan, and Ion Stoica. 2017. "RLlib: Abstractions for Distributed Reinforcement Learning," December.

Metz, Luke, Niru Maheswaranathan, Jonathon Shlens, Jascha Sohl-Dickstein, and Ekin D Cubuk. 2019. "Using Learned Optimizers to Make Models Robust to Input Noise," June.

Miller, E G, N E Matsakis, and P A Viola. 2000. "Learning from One Example Through Shared Densities on Transforms." In *Proceedings IEEE Conference on Computer Vision and Pattern Recognition. CVPR 2000 (Cat. No.PR00662)*, 1:464–71 vol.1.

Nagabandi, Anusha, Ignasi Clavera, Simin Liu, Ronald S Fearing, Pieter Abbeel, Sergey Levine, and Chelsea Finn. 2018. "Learning to Adapt in Dynamic, Real-World Environments Through Meta-Reinforcement Learning," March.

Oreshkin, Boris, Pau Rodríguez López, and Alexandre Lacoste. 2018. "TADAM: Task Dependent Adaptive Metric for Improved Few-Shot Learning." In *Advances in Neural Information Processing Systems 31*, edited by S Bengio, H Wallach, H Larochelle, K Grauman, N Cesa-Bianchi, and R Garnett, 721–31. Curran Associates, Inc.

Park, Eunbyung, and Junier B Oliva. 2019. "Meta-Curvature," February.

Paszke, Adam, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, et al. 2019. "PyTorch: An Imperative Style, High-Performance Deep Learning Library." In *Advances in Neural Information Processing Systems 32*, edited by H Wallach, H Larochelle, A Beygelzimer, F d'Alché-Buc, E Fox, and R Garnett, 8026–37. Curran Associates, Inc.

Raghu, Aniruddh, Maithra Raghu, Samy Bengio, and Oriol Vinyals. 2019. "Rapid Learning or Feature Reuse? Towards Understanding the Effectiveness of MAML," September.

Rothfuss, Jonas, Dennis Lee, Ignasi Clavera, Tamim Asfour, and Pieter Abbeel. 2018. "ProMP: Proximal Meta-Policy Search."

Sutton, Richard S. 1992. "Adapting Bias by Gradient Descent: An Incremental Version of Delta-Bar-Delta," June.

Vinyals, Oriol, Charles Blundell, Timothy Lillicrap, Koray Kavukcuoglu, and Daan Wierstra. 2016. "Matching Networks for One Shot Learning." In *Advances in Neural Information Processing Systems 29*, edited by D D Lee, M Sugiyama, U V Luxburg, I Guyon, and R Garnett, 3630–8. Curran Associates, Inc.

Xu, Zhongwen, Hado van Hasselt, and David Silver. 2018. "Meta-Gradient Reinforcement Learning," May.